

matrices and arrays in R

A matrix is a vector with two additional attributes: the number of rows and the number of columns. Matrices are also vectors; they have modes (classes). However, vectors are not one-rows/column matrices.

Matrices are special cases of R type object: arrays. Unlike matrices, arrays can be multi-dimensional; having rows, columns, and layers. Matrix row and column subscripts (indexes) begin with 1. The internal storage of a matrix is column-major order, meaning all of column 1 is stored initially, then all of column 2, and so on.

However, it can be specified as an argument to build a matrix with row-major order, opposed to the default:

```
1 > y <- matrix(c(1, 2, 3, 4), nrow = 2, ncol = 2)      #build 2x2 matrix
2 > y
3      [,1] [,2]
4 [1,]    1    3
5 [2,]    2    4
6
7 > y <- matrix(1:4 nrow = 2)                          #only nrow or ncol needed
8 > y
9      [,1] [,2]
10 [1,]    1    3
11 [2,]    2    4
12
13 > y <- matrix(nrow = 2, ncol = 2)                   #individually specifying elements
14 > y
15 > y[1, 1] <- 1
16 > y[2, 1] <- 2
17 > y[3, 1] <- 3
18 > y[1, 2] <- 3
19 > y[2, 2] <- 4
20      [,1] [,2]
21 [1,]    1    3
22 [2,]    2    4
23
24 > m <- matrix(c(1, 2, 3, 4, 5, 6), nrow = 2, byrow = TRUE) #specify byrow argument
25 > m
26      [,1] [,2] [,3]
27 [1,]    1    2    3
28 [2,]    4    5    6
```

Common Linear Algebra operations can be performed on matrices, such as matrix multiplication, matrix scalar multiplication, matrix addition, matrix subtraction, etc.

```
1 > y %% y      #mathematical matrix multiplication
2      [,1] [,2]
3 [1,]    7   15
4 [2,]   10   22
5
6 > 3*y         #mathematical multiplication of matrix by scalar
7      [,1] [,2]
8 [1,]    3    9
9 [2,]    6   12
10
11 > y+y        #mathematical matrix addition
12      [,1] [,2]
13 [1,]    2    6
14 [2,]    4    8
```

matrix subsetting `[]` in R

Matrices can be subset like vectors, with the addition of column arguments.

```
1 > z <- matrix(c(1,2,3,4,1,1,0,0,1,0,1,0), ncol = 3) #build 3x4 matrix z
2 > z
3      [,1] [,2] [,3]
4 [1,]    1    1    1
5 [2,]    2    1    0
6 [3,]    3    0    1
7 [4,]    4    0    0
8
9 > z[,2:3] #subset columns 2-3 of z
10      [,1] [,2]
11 [1,]    1    1
12 [2,]    1    0
13 [3,]    0    1
14 [4,]    0    0
15
16 > z[1:2,1] #subset rows 1-2 and column 1 of z
17 [1] 1 2
18
19 > z[,-1] #subset the complement of z, column 1
20      [,1] [,2]
21 [1,]    1    1
22 [2,]    1    0
23 [3,]    0    1
24 [4,]    0    0
25
26 > z[1,] <- matrix(c(7,7,7), nrow = 3) #assign 1st row with vector
27 > z
28      [,1] [,2] [,3]
29 [1,]    7    7    7
30 [2,]    2    1    0
31 [3,]    1    0    1
32 [4,]    4    0    0
33
34 > x <- matrix(nrow = 5, ncol = 4) #assign matrix z to rows 2-5 and columns 2-4 of x
35 > x[2:5,2:4] <- y
36 > x
37      [,1] [,2] [,3] [,4]
38 [1,]   NA   NA   NA   NA
39 [2,]   NA    7    7    7
40 [3,]   NA    2    1    0
41 [4,]   NA    1    0    1
42 [5,]   NA    4    0    0
```

matrix filtering in R

Matrices can be filtered similar to vectors; however, it is important to remain mindful of the applied syntax.

```
1 > x <- matrix(c(1,2,2,3,3,4), nrow = 3, byrow = TRUE)
2 > x
3      [,1] [,2]
4 [1,]    1    2
5 [2,]    2    3
6 [3,]    3    4
7
8 > x[x[,2] >= 3,]
9      [,1] [,2]
10 [1,]    2    3
11 [2,]    3    4
12
13 > j <- x[,2] >= 3
14 > j
15 [1] FALSE  TRUE  TRUE
16
17 > x[j,]
18      [,1] [,2]
19 [1,]    2    3
20 [2,]    3    4
```

The assignment of variable **j** to filter matrix **x** is a vectorized operation; all of the below hold true:

- ... The object `x[, 2]` is a vector
- ... The operator `>=` compares two vectors
- ... The number 3 was recycled to a vector of 3's

Similarly, a variable can be used outside of the variable that the filtering is applied:

```
1 > m <- matrix(1:6, nrow = 3)    #build a 3x2 matrix m
2 > m
3      [,1] [,2]
4 [1,]    1    4
5 [2,]    2    5
6 [3,]    3    6
7
8 > m[m[,1] > 1 & m[,2] > 5,]    #filter matrix m to the row(s) where the first column
9 [1] 3 6                        values > 1 and the second column values > 5
```

Note that the above operation should have returned a 1x2 matrix, as opposed to a two-element vector. This is caused by the incorrect data type being assigned. The solution to avoiding this outcome is to apply the `drop` argument in the function; this will tell R to retain the two-dimensional nature of the data.

Additionally, since matrices are vectors, vector operations can be applied appropriately:

```
1 > m
2      [,1] [,2]
3 [1,]    1    4
4 [2,]    2    5
5 [3,]    3    6
6
7 > which(m > 2)                #returns a 4-element vector of all matrix elements of m > 2
8 [1] 3 4 5 6
```

applying functions to matrix rows and columns $f(x)$ in R

The `apply()` functions are among the most used functions in R; `apply()`, `tapply()`, `lapply()`. Each instructs R to call user-defined functions against each row or each column in a matrix.

A generalized formation of the `apply()` function:

```
1 > apply(X, margin, FUN, ...)
```

The arguments in the above `apply()` function form are as follows:

- ... **X** is the matrix
- ... **margin** is the dimension; **1** applies the function to **rows**, **2** applies the function to **columns**
- ... **FUN** is the applied function
- are optional arguments to be supplied to function **f**

Below applies the `mean()` function matrix **z** as an example:

```
1 > z <- matrix(1:6, nrow = 3)      #build 3x2 matrix z
2 > z
3      [,1] [,2]
4 [1,]    1    4
5 [2,]    2    5
6 [3,]    3    6
7
8 > apply(z,2,mean)                  #apply the mean function to matrix z by columns
9 [1] 2 5
10
11 > colMeans(z, 2L)                 #apply the colMeans function to matrix z by columns
12 [1] 2 5
```

Note that the `colMeans()` function performs the same operation and could be used instead; however, the above example is for illustrative purposes of the `apply()` function.

```
1 > f <- function(x) x/c(2,8)        #build function f that divides the elements
2 > y <- apply(z,1,f)                 of x by the two-dimensional vector(2, 8)
3 > y
4      [,1] [,2] [,3]
5 [1,]  0.5 1.000 1.50
6 [2,]  0.5 0.625 0.75
7
8 > t(y)                              #transpose matrix y to reflect the
9      [,1] [,2]                      dimensions of the original matrix z
10 [1,]  0.5 0.500
11 [2,]  1.0 0.625
12 [3,]  1.5 0.750
```

adding and deleting matrix rows and columns [x] in R

Matrices are fixed in terms of rows and columns; thus, the latter cannot be technically *deleted*. Instead, matrices can be **assigned** and **reassigned** resulting in the same net effect.

Revisiting the reassignment of **vectors** to change size:

```
1 > x <- c(12, 5, 13, 16, 8) #build five-element vector
2 > x
3 [1] 12  5 13 16  8
4
5 > x <- c(x, 20) #append value 20
6 > x
7 [1] 12  5 13 16  8 20
8
9 > x <- c(x[1:3], 20, x[4:6]) #insert value 20
10 > x
11 [1] 12  5 13 20 16  8 20
12
13 > x <- x[-2:-4] #delete elements 2 through 4
14 > x
15 [1] 12 16  8 20
```

Note: reassignment occurs often unseen. For example, even the assignment **x[2] <- 12** is a reassignment.

Analogous operations can be used to change the size of a **matrix**:

```
1 > one #build four-element vector "one"
2 [1] 1 1 1 1
3
4 > z <- matrix(c(1:4,1,1,0,0,1,0,1,0), nrow = 4) #build 4x3 matrix z
5 > z
6      [,1] [,2] [,3]
7 [1,]    1    1    1
8 [2,]    2    1    0
9 [3,]    3    0    1
10 [4,]    4    0    0
11
12 > cbind(one,z) #bind vector "one" to matrix z
13      one
14 [1,]    1 1 1 1
15 [2,]    1 2 1 0
16 [3,]    1 3 0 1
17 [4,]    1 4 0 0
18
19 > cbind(1,z) #bind values 1 to matrix z by recycling
20      [,1] [,2] [,3] [,4]
21 [1,]    1    1    1    1
22 [2,]    1    2    1    0
23 [3,]    1    3    0    1
24 [4,]    1    4    0    0
```

```

1  > m <- matrix(1:6, nrow = 3)           #build 3x2 matrix m
2  > m
3      [,1] [,2]
4  [1,]    1    4
5  [2,]    2    5
6  [3,]    3    6
7
8  > m <- m[c(1,3),]                       #delete row 2 through reassignment
9  > m
10     [,1] [,2]
11  [1,]    1    4
12  [2,]    3    6

```

matrix and vector distinction ■ ■ in R

Considering that matrices are also vectors:

```

1  z <- matrix(1:8, nrow = 4)             #build 4x2 matrix z
2  > z
3      [,1] [,2]
4  [1,]    1    5
5  [2,]    2    6
6  [3,]    3    7
7  [4,]    4    8

```

As a vector, the length of matrix z can be queried:

```

1  > length(z)
2  [1] 8

```

However, a matrix is in its own class as a "matrix" object (an S3 class).

```

1  > class(z)                               #print the class of matrix z
2  [1] "matrix"
3
4  > attributes(z)                          #print the attributes of matrix z
5  $dim
6  [1] 4 2
7
8  > dim(z)                                 #the dimensions can also be printed
9  [1] 4 2
10
11 > nrow(z)                                #print the number of matrix z rows
12 [1] 4
13
14 > ncol(z)                                #print the number of matrix z columns
15
16 [1] 2

```

Lastly, the objects themselves can be printed by calling their names:

```

1  > nrow
2  function (x)
3  dim(x)[1L]
4  <bytecode: 0x00000000095e36c8>
5  <environment: namespace:base>

```

The benefits of object-oriented programming are shown with functions whose arguments are a matrix, or matrices. The number of rows/columns are available rather than needing to be supplied in arguments.

avoiding dimension reduction in R

There are many cases in statistics where dimension reduction is the goal. However, the latter is only the case when the result of losing dimensions is the *intent*.

```
1 > z                                     #build 4x2 matrix z
2      [,1] [,2]
3 [1,]    1    5
4 [2,]    2    6
5 [3,]    3    7
6 [4,]    4    8
7
8 > r <- z[2,]                          #subset matrix z for row 2; the result is a vector, not a matrix
9 > r
10 [1] 2 6
```

Noting that the result of the above operation is a vector, the original matrix class is lost. Concretely, the result is a two-element vector, as opposed to a 1x2 matrix. The latter is proven as follows:

```
1 > attributes(z)                       #matrix z is a 4x2 matrix
2 $dim
3 [1] 4 2
4
5 > attributes(r)                       #vector r has no discernable attributes
6 NULL
7
8 > str(z)                              #matrix z has 2 indices for rows and columns
9 int [1:4, 1:2] 1 2 3 4 5 6 7 8
10
11 > str(r)                              #vector r has a single index of range 1:2
12 int [1:2] 2 6
```

The problems caused by unintentional dimension reductions within R code can result in general purpose code performing as expected in normal conditions; but failing in special cases. For example, assume a submatrix is extracted from a given matrix, followed by matrix operations on the submatrix. If the submatrix only has a single row, R coerces to a vector and fails in the following matrix operations.

Unintentional dimensionality reduction can be avoided through the use of the **drop** argument.

```
1 > r <- z[2,, drop = FALSE]            #maintain matrix integrity with the drop() argument
2 > r
3      [,1] [,2]
4 [1,]    2    6
5
6 > dim(r)                              #the dimensions of matrix r now consists of 2 columns
7 [1] 1 2
```

The consideration of **drop** as an argument is due to “[” actually serving as a function, like “+” noted prior.

```
1 > z[3,2]                             #subset matrix z to return row 3 and column 2
2 [1] 7
3
4 > "["(z,3,2)                          #subset matrix with the "[" function z to return row 3 and column 2
5 [1] 7
```

Additionally, the **as.matrix()** function can be used to treat an existing vector as a matrix.

```
1 > u <- c(1,2,3)           #construct three-element vector u
2 > u
3 [1] 1 2 3
4
5 > attributes(u)           #vector u contains no discernible attributes, as expected
6 NULL
7
8 > v <- as.matrix(u)       #construct matrix v by treating vector u as a matrix
9 > v
10      [,1]
11 [1,]    1
12 [2,]    2
13 [3,]    3
14
15 > attributes(v)          #matrix v contains dimensions of 3 rows and 1 column
16 $dim
17 [1] 3 1
```

naming matrix rows and columns 🏠 in R

Rows and columns in matrices are often referred to by their corresponding numbers. However, matrix rows and numbers can be assigned names with the **rownames()** and **colnames()** functions as follows:

```
1 > z                       #print 2x2 matrix z
2      [,1] [,2]
3 [1,]    1    3
4 [2,]    2    4
5
6 > colnames(z)             #matrix z has no names assigned to columns
7 NULL
8
9 > colnames(z) <- c("a", "b") #assign names to columns of matrix z
10 > z
11      a b
12 [1,] 1 3
13 [2,] 2 4
14
15 > colnames(z)             #matrix z has 2 columns names "a" and "b"
16 [1] "a" "b"
17
18 > z[, "a"]               #column names can be used as references
19 [1] 1 2
```


higher-dimensional arrays in R

A typical matrix in R consists of rows and corresponding observations; a two-dimensional data structure. Assume the hypothetical where the same observations occur at different time frames amongst the sample (students taking tests in multiple periods during a school term). Time becomes the third dimension in R and these datasets are called **arrays**. The following example illustrates a series of two tests students take during a given period:

```
1 > firsttest <- matrix(c(46,21,50,30,25,48), ncol = 2)      #construct fist test matrix
2 > secondtest <- matrix(c(46,41,50,43,35,49), nrow = 3)     #construct fist test matrix
3
4 > firsttest
5      [,1] [,2]
6 [1,]   46  30
7 [2,]   21  25
8 [3,]   50  48
9
10 > secondtest
11     [,1] [,2]
12 [1,]   46  43
13 [2,]   41  35
14 [3,]   50  49
```

The **array()** function is defined above to construct a two-layer array, each consisting of three rows and two columns. The number of layers are assigned with the second '**2**' in the **dim = c(3,2,2)** argument.

```
1 > attributes(tests)          #the attributes of the tests array are printed
2 $dim
3 [1] 3 2 2
4
5 > tests[3,2,1]               #the score on the second portion of test 1 for student 3
6 [1] 48
7
8 >                             #merge firsttest and secondtest matrices into an array
9 >
10 > tests <- array(c(firsttest, secondtest), dim = c(3,2,2))
11 > tests
12 , , 1
13
14     [,1] [,2]
15 [1,]   46  30
16 [2,]   21  25
17 [3,]   50  48
18
19 , , 2
20
21     [,1] [,2]
22 [1,]   46  43
23 [2,]   41  35
24 [3,]   50  49
```

In addition to building a three-dimensional array by combining the two matrices above, four-dimensional matrices can be constructed by combining two, or more, three-dimensional arrays, and so on.

A common use of arrays is applied to calculating tables, discussed later.